Symbolic Execution in CPAchecker

Thomas Lemberger
lemberge@fim.uni-passau.de

ABSTRACT

The value analysis of CPAchecker, a successful tool for configurable software verification, possesses high performance but is not able to handle non-deterministic values in general. This results in a high amount of false alarms. While counterexample checks with other analyses are currently used to mitigate this problem, we extend the existing value analysis to allow the handling of non-deterministic values and introduce a new configurable program analysis (CPA) to track relations between such values. Our evaluation shows that this allows analyzing programs that use non-deterministic values correctly without the need for counterexample checks.

1. MOTIVATION

With the rise of ubiquitious computing, reactive software systems that constantly receive input from the outside become more and more common. Because of their non-functional behaviour, the unreliability of testing with a finite number of test values becomes even more severe for such programs. In constrast to this, automatic software verification offers a more reliable alternative by analyzing all possible behaviours of a program. Configurable software verification[5] and one of its implementations CPAchecker[6] represent a recent approach to software verification that has been successful in multiple iterations of the Competition on Software Verification (SV-COMP) [1] [2] [3].

A prominent CPA used in CPAchecker is the ValueAnalysisCPA [7]. While it offers high efficiency, it only tracks explicit values and cannot handle non-deterministic ones. Consider the simple example program in Listing 1. Its analysis by the ValueAnalysisCPA is displayed in Figure 1. It can easily be seen that the non-deterministic value of a cannot be handled and as such is discarded. Because of this imprecision, necessary information about the relation between a and b gets lost and the safety property in Line 14 is seen as violated.

This characteristic results in a high amount of false positives. In practice, this problem is countered by the use of coun-

```
1  extern __nondet_int ();
2
3  int main () {
4      int a = __nondet_int ();
5      int b;
6
7      if (a >= 0) {
8          b = a;
9
10     } else {
11         b = a + 1;
12     }
13
14     if (b < a) {
15 ERROR:
16         return −1;
17     }
18 }
```

**Listing 1: A simple non-deterministic program**

terexample checks and strengthening through other CPAs. Instead of that, we introduce symbolic values to the existing ValueAnalysisCPA to allow tracking of non-deterministic values. In addition, we specify a new ConstraintsCPA to track constraints on these symbolic values.

First, we will give the formal specification of a composite CPA (which we will call SymbolicExecutionCPA) using these two components. After this, an implementation of this CPA based on CPAchecker is presented. An evaluation on benchmarks taken from SV-COMP 2015 will show that the addition of symbolic value tracking allows the correct analysis of programs that could priorly not be solved by the ValueAnalysisCPAalone.

2. SPECIFICATION

The SymbolicExecutionCPA is a composite CPA of the *LocationCPA*, which tracks the current location in the syntactic structure of a program, the *SymbolicValueCPA*, which tracks and computes the deterministic and non-deterministic values of variables and the *ConstraintsCPA*, which tracks encountered assumptions in form of constraints for each location's variables. The formal definitions of CompositeCPA and LocationCPA are used as defined in [5].

For the sake of simplicity, we assume that all values are

Figure 1: Graph illustrating the analysis of the program in Listing 1 by the ValueAnalysisCPA. The abstract state of the ValueAnalysisCPA is always empty as it does not track non-deterministic values. Because of this, states are merged at Location 6.

integers and that a program only consists of assignments and assumptions.

2.1 SymbolicValueCPA

The described CPA is an extension of the priorly mentioned ValueAnalysisCPA. Instead of using a bottom element $\perp$, we represent an unreachable state through the absence of a valid transfer. The SymbolicValueCPA is defined as

$$\mathbb{S} = (D_{\mathbb{S}}, \rightsquigarrow_{\mathbb{S}}, merge^{sep}\ stop^{sep})$$

with abstract domain $D_{\mathbb{S}}$, transfer relation $\rightsquigarrow_{\mathbb{S}}$, merge operator $merge^{sep}$ and stop operator $stop^{sep}$.

The abstract domain $D_{\mathbb{S}}$ is defined as $D_{\mathbb{S}} = (C, \mathcal{E}, \llbracket \cdot \rrbracket)$ with C being the set of concrete program states, $\mathcal{E}$ the semi-lattice of possible abstract states and $\llbracket \cdot \rrbracket$ the concretization function.

$$\mathcal{E} = (V_{\mathbb{S}}, \sqsubseteq, \sqcup, v_{\top})$$

The elements of the semi-lattice are partial functions of $V_{\mathbb{S}} = X \rightarrow (\mathbb{Z} \cup \mathcal{Z}_{\mathbb{S}})$ mapping program variables in its definition range to concrete values of $\mathbb{Z}$ or to symbolic values $\mathcal{Z}_{\mathbb{S}} = S_I \cup S_E$. S_I describes all symbolic identifiers and S_E is the set of all valid symbolic expressions. Each expression is a symbolic expression if at least one symbolic identifier occurs in it. The definition range of a function f is defined as $def(f) = \{x \mid \exists y : f(x) = y\}$. The definition range of an abstract variable assignment of the type $V_{\mathbb{S}}$ consists of all program variables $x \in X$ whose value is known, with X being the set of all program variables occurring in the program. If a program variable $x \in X$ has no assignment for an abstract state $v \in V_{\mathbb{S}}$, its value is unknown. There are multiple reasons for unknown values: Uninitialized variables are unknown, but input values or calls to unknown functions and operations are, too.

To define $\sqsubseteq$, we first have to define the relation between concrete and abstract states. A concrete state $c \in C$ is

compliant to an abstract variable assignment v, if an arbitrary, but valid assignment of concrete values to symbolic identifiers $d : S_I \rightarrow \mathbb{Z}$ exists, so that for all $x \in def(v)$ one of the following conditions is fulfilled: (1) the abstract assignment is a concrete value that equals the value in the concrete state, that means $c(x) = v(x)$, or (2) the abstract assignment $v(x)$ is a symbolic value that can be evaluated to $c(x)$ by replacing all occurring symbolic identifiers $i \in S_I$ with $d(i)$. The concretization function $\llbracket \cdot \rrbracket$ then assigns all compliant concrete states to an abstract state v.

For two abstract states $v, v' \in V_{\mathbb{S}}$, the concrete states that are covered by v are a subset of the ones covered by v', noted as $v \sqsubseteq v'$, if all of the following conditions hold: (1) v' must only contain value assignments also present in v, that is $def(v') \subseteq def(v)$, (2) $\forall x \in def(v') : v'(x) \in \mathbb{Z} \Rightarrow v'(x) = v(x)$ and (3) a bijective function alias $: S_I \rightarrow S_I$ exists that maps each symbolic identifier of S_I to another symbolic identifier, so that $\forall x \in def(v') : v'(x) \in \mathcal{Z}_{\mathbb{S}} \Rightarrow v(x)$ results from $v'(x)$ by replacing all $i \in S_I$ occurring in $v'(x)$ with alias(i). Constraint (3) ensures that not the symbolic values themselves are compared, but the number of concrete states they represent. Consider $v = \{(a, s1), (b, s1+5)\}$ and $v' = \{(a, s2), (b, s2+5)\}$ with $a, b \in X$ and $s1, s2 \in S_I$. As $s1$ and $s2$ represent any possible value, v represents the same amount of concrete states as v', that is all $c \in C$ with $c(b) = c(a) + 5$ and an arbitrary $c(a)$. Because of this, $v' \sqsubseteq v$ holds (and in this case even $v \sqsubseteq v'$).

Note that $v' \sqsubseteq v \Rightarrow \llbracket v' \rrbracket \subseteq \llbracket v \rrbracket$, but $\llbracket v' \rrbracket \subseteq \llbracket v \rrbracket \not\Rightarrow v' \sqsubseteq v$.

The join $\sqcup$ represents the least upper bound of two lattice elements $v, v' \in V_{\mathbb{S}}$ with

$$(v \sqcup v')(x) = v(x) \text{ for all } x \in def(v) \cap def(v') \text{ and } v(x) = v'(x).$$

The top element of $\mathcal{E}$ is defined as $v_{\top} = \{\}$, i.e. no known assignments exist. It represents all possible concrete assignments, formally expressed as $\llbracket v_{\top} \rrbracket = C$.

The transfer relation $\rightsquigarrow_{\mathbb{S}}$ contains the transfer $v \overset{g}{\rightsquigarrow}_{\mathbb{S}} v''$, if one of the following conditions is true:

1. $g = (l, assume(p), l'), \phi(p, v)$ is satisfiable and:

$$v''(x) = \begin{cases} c & \text{if } c \text{ only satisfying assignment} \\ & \text{for } \phi(p, v) \text{ and } x \notin def(v) \\ y & \text{if } x \notin def(v) \text{ and } x \text{ appears in} \\ & p.\ y \in S_I \text{ and } y \text{ is a new value} \\ & \text{that has not been used in any} \\ & \text{other state before} \\ v(x) & \text{otherwise} \end{cases}$$

with

$$\phi(p, v) := p \wedge (\bigwedge_{\substack{x \in def(v), \\ v(x) \in \mathbb{Z}}} x = v(x))$$

Note: ϕ performs an over-approximation, as only variables $x \in def(v)$ with a concrete value are considered.

2. $g = (l, w := e, l')$ and:

$$v''(x) = \begin{cases} \text{eval}(e, v') & \text{if } x = w \\ v'(x) & \text{if } x \in def(v) \text{ and } x \neq w \end{cases}$$

with

$$v'(x) = \begin{cases} y & \text{if } x \notin \text{def}(v) \text{ and } x \text{ appears in} \\ & e. \ y \in S_I \text{ and } y \text{ is a new value} \\ & \text{that has not been used in any} \\ & \text{other state before} \\ v(x) & \text{if } x \in \text{def}(v) \end{cases}$$

and $\text{eval}(e, v')$ defined as the evaluation of an expression e with the given assignment. If any symbolic value occurs in $v(e)$, it is only partially evaluated. In this case, $\text{eval}(e, v') \in \mathcal{Z}_{\mathbb{S}}$.

3. $v'' = v_\top$.

The merge operator $merge^{sep}$ is defined as $merge^{sep}(e, e') = e'$. That means no merge is performed for different value states at the same location. The stop operator $stop^{sep}$ considers every state of the reached set independently when checking for coverage. It is defined as $stop^{sep}(e, R) = \exists e' \in R : e \sqsubseteq e'$.

2.2 ConstraintsCPA

The ConstraintsCPA is a CPA

$$\mathbb{C} = (D_{\mathbb{C}}, \leadsto_{\mathbb{C}}, merge^{sep}, stop^{sep})$$

that tracks constraints (i.e. boolean formulas) on variables that are created by assume edges. The abstract domain $D_{\mathbb{C}}$ is defined by

$$D_{\mathbb{C}} = (C, \mathcal{C}, \llbracket \cdot \rrbracket)$$

with concrete states C, the semi-lattice $\mathcal{C}$ and concretization function $\llbracket \cdot \rrbracket$.

The abstract states described by $\mathcal{C} = (2^\gamma, \sqsubseteq, \sqcup, \top)$ are subsets of γ. γ is the set of all possible boolean expressions over $\mathbb{Z} \cup \mathcal{Z}_{\mathbb{S}}$, including symbolic expressions. An abstract state $a \in \mathcal{C}$ can be interpreted as the conjunction of all its constraints.

For two given states $a, a' \in \mathcal{C}$, state a is less than or equal to a', if a bijective funtion $\text{alias} : S_I \to S_I$ exists, so that $a'' \subseteq a$ with a'' resulting from a' by replacing all symbolic identifiers $i \in S_I$ occurring in constraints of a' with $\text{alias}(i)$.

The concretization function $\llbracket \cdot \rrbracket$ maps an abstract state to all concrete states that satisfy this abstract state's constraints.

$$\llbracket a \rrbracket = \{ c \in C | \ c \vDash \bigwedge_{\varphi \in a} \varphi \}$$

Note: We assume that the empty conjunction represents true, that means $\bigwedge_{\varphi \in \phi} \varphi = true$ for $\phi = \{\}$.

The transfer relation $\leadsto_{\mathbb{C}}$ contains the transfer $a \xrightarrow{g}_{\mathbb{C}} a'$ if one of the following conditions is fulfilled: (1) $g = (l, assume(p), l')$, $a' = a \cup p$ and a does not contain any variable $x \in X$ (by specifying that a must not contain any variable we enforce that variables are always replaced by concrete or symbolic values through strengthening), or (2) $g = (l, w := e, l')$ and $a' = a$.

As p is only added as a constraint if no variable of X occurs in a, all resulting states of this transfer relation will be

satisfiable, unless p is a contradiction (e.g. $a \neq a$). Because of this, we completely disregard checking for satisfiability of constraints here. Instead, we rely on satisfiability checks during strengthening by other CPAs. In the following we will use the priorly defined SymbolicValueCPA to replace newly added constraints' variables with values and check for the whole state's satisfiability.

2.3 Composition of CPAs

$\mathcal{C}_{\mathbb{LSC}} = (\mathbb{L}, \mathbb{S}, \mathbb{C}, \leadsto_x, merge^{sep}, stop^{sep})$ is the composite CPA of the LocationCPA $\mathbb{L}$, the SymbolicValueCPA $\mathbb{S}$ and the ConstraintsCPA $\mathbb{C}$. The transfer relation

$$\leadsto_x : D_{\mathbb{L}} \times D_{\mathbb{S}} \times D_{\mathbb{C}} \to D_{\mathbb{L}} \times D_{\mathbb{S}} \times D_{\mathbb{C}}$$

contains the transfer $(l, v, a) \xrightarrow{g}_x (l', v', a')$ if $l \xrightarrow{g}_{\mathbb{L}} l'$, $v \xrightarrow{g}_{\mathbb{S}} v'$ and $a \xrightarrow{g}_{\mathbb{C}} a'$ and if the strengthen operator $\downarrow_{\mathbb{C}, \mathbb{S}}$ is defined for a' and v'.

The strengthen operator $\downarrow_{\mathbb{C}, \mathbb{S}} : \mathbb{C} \times \mathbb{S} \to \mathbb{C}$ uses the value assignment $v(x)$ of the given second operand to replace variables of the latest, unstrengthened constraint of the given constraints state. This way, meaningful constraints are created that show relations between symbolic values and can be checked for satisfiability.

$\downarrow_{\mathbb{C}, \mathbb{S}} (a, v) = a'$ is defined if a' is the result of replacing all program variables $x \in X$ occurring in a with $v(x)$ and if $\bigwedge_{\varphi \in a'} \varphi$ is satisfiable.

An analysis of the example program (c.f. Listing 1) using symbolic execution can be seen in Figure 2. By storing information about the relationship of variables a and b, the CPA can derive at location 6/6' that b cannot be smaller than a.

3. EVALUATION

The specified CPA was implemented in CPAchecker[6] using the existing CompositeCPA and LocationCPA. The existing ValueAnalysisCPA was extended to fulfil the specification of the SymbolicValueCPA, while the constraints CPA was implemented as a new, individual CPA. To check the satisfiability of constraints, an external SAT solver is used. For these benchmarks, MathSAT5 is used with the bitvector theory, encoding float values as floats. The analysis is part of CPAchecker's trunk since revision 16052 and can be activated by using the configuration `valueAnalysis-symbolic`. Revision 16433 was chosen for benchmarks.

Benchmarks were performed on a subset of the SV-COMP 2015 test set, excluding sets CPAchecker's ValueAnalysis-CPA has no support for. These excluded sets are "Concurrency", "Memory Safety", "Recursive" and "Termination". An overview of all test sets can be found at [4]. Tests were executed on Intel Xeon E7-4870 machines at 2.40GHz with 80 cores and 322GB of available memory, with a memory limit of 15.0GB, a time limit of 900 seconds and a core limit of two CPUs per run. A Java heap memory limit of 10.5GB per run was used.

All benchmark results are available at `http://leostrakosch.github.io/symbolicValueAnalysis`.
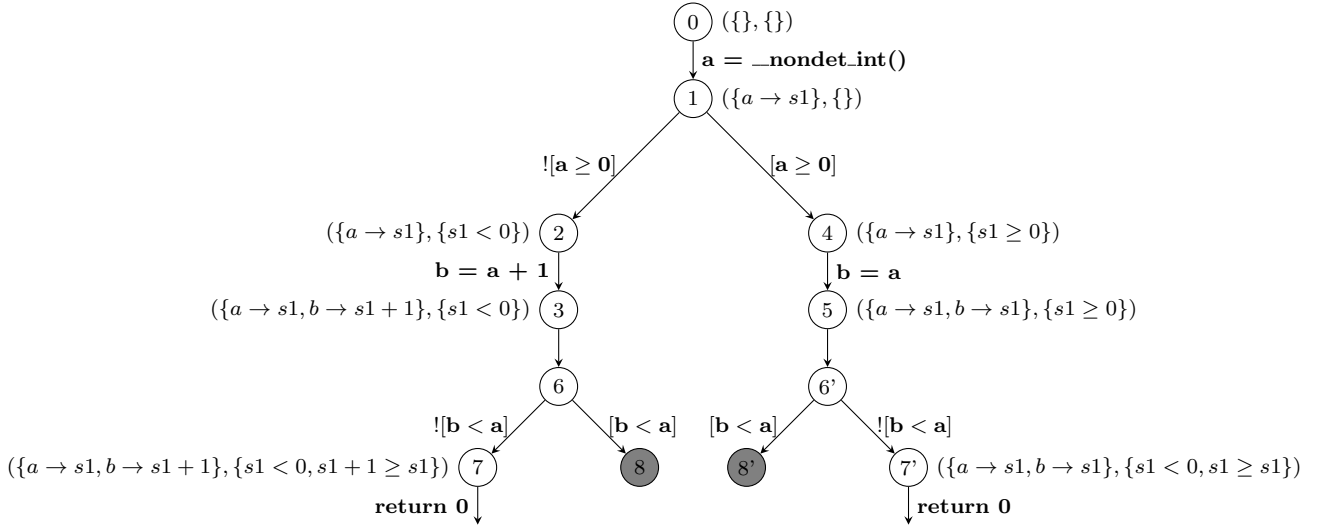
Figure 2: Analysis of the program in Listing 1 by the SymbolicExecutionCPA. The locations reported as unreachable are tinted grey. The state is provided as a tuple at each relevant node, the first element being the symbolic state of $\mathbb{S}$, the second the constraints of $\mathbb{C}$. The state of the LocationCPA can be derived from the node's numbers.

	Value	SymEx	Overall
correct	1995	833	3038
true positives	632	246	802
true negatives	1363	587	2236
unique true positives	393	7	-
unique true negatives	804	28	-
false positives	271	52	-
unique false positives	222	3	-
false negatives	0	0	-
timeouts	662	1970	-
errors	110	182	-
memory errors	0	1	-

Figure 3: Result of runs of value analysis (Value) and symbolic execution (SymEx) on SV-COMP 2015 test sets. "true positive" represents a found property violation in a program containing a property violation, "true negative" represents no found violation in a program without violations.

Figure 3 shows the performance of the default ValueAnalysisCPA without counterexample-guided abstraction refinement (CEGAR [7]) or counterexample checks next to the performance of the new symbolic execution configuration. Generally speaking, basic value analysis strongly outperforms symbolic execution in numbers due to symbolic execution taking too long. Lots of timeouts occur due to (a) path explosion, an exponential increase in possible paths based on the number of branches (that is if-statements and loops) as a result of the low level of abstraction of symbolic execution, and (b) the bad performance of SAT checks with a large number of variables and arithmetic operations like non-linear multiplication. Path explosion can easily be illustrated when looking at the algorithm statistics of a simple program using multiple if- and goto-statements.[1] While ba-

---
[1]Namely program `ldv-validator-v0.6/linux-`

sic value analysis only needs 1241 iterations to find a possible property violation, with the biggest waitlist consisting of 42 entries, analysis with symbolic execution needs 14572 iterations with the biggest waitlist consisting of 543 states to find the same property violation. The scatter plot in Figure 4 displays this problem. For some tasks, symbolic execution is able to compute the result with less successors than value analysis due to its stronger precision. But in many cases, symbolic execution computes far more successors because of this. In addition to this enormous growth in iterations, SAT
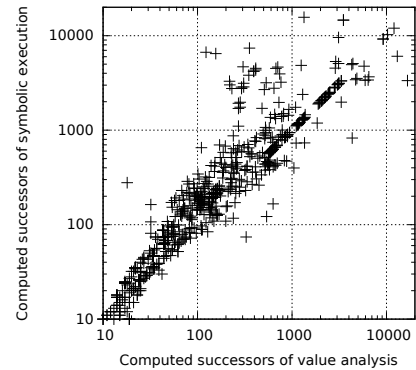

Figure 4: Number of computed successors for ValueAnalysisCPA and SymbolicExecutionCPA. Only the computed successors for benchmark runs are displayed which returned a result for both SymbolicExecutionCPA and ValueAnalysisCPA and for which the result was the same. The upper bound used for computed successors is 20,000.

checks are responsible for up to 95.0% of CPU time in analyses of programs using non-deterministic values that are *not*

`stable-9ec4f65-1-110_1a-drivers-rtc-rtc-tegra.ko-entry_point_false-unreach-call.cil.out.c`

specifically constructed for challenging SAT solvers. This influences the computation time per successor. As can be seen in Figure 5, there are many cases in which SymbolicExecutionCPA takes significantly longer to compute successors than ValueAnalysisCPA.

Figure 5: Number of computed successors per second for ValueAnalysisCPA and SymbolicExecution-CPA

Despite this reduced speed, symbolic execution's ability to track non-deterministic values allows it to get a correct result in 35 cases in which value analysis does not. In seven of these value analysis exceeds the timelimit and in 28 value analysis provides a false positive. Symbolic execution does never produce a false result when value analysis produces a correct one.

When using a timelimit of 1800 seconds, symbolic execution is able to compute 9 more results than with a timelimit of 900 seconds, of which seven are correct.

Of these long-taking runs, three are the result of path explosion, while the other five only consist of few iterations, but contain expressions resulting in long-taking SAT checks, namely bitvector computations, non-deterministic floats and non-linear multiplications.

These performance issues can be improved by using another theory in SAT checks. By using an integer theory instead of bitvector, runtime can be greatly improved in some cases. The scatter plot in Figure 7 illustrates the runtime improvements of integer theory over bitvector. This results in a lower precision and an unsound analysis, though. Figure 6 shows the difference between the use of a bitvector theory with floats and the use of an integer theory.

4. CONCLUSION

By extending the existing ValueAnalysisCPA with the ability to track non-deterministic values and the introduction of the new ConstraintsCPA, we were able to reduce the number of false positives of value analysis without the need for counterexample checks. In contrast to predicate analysis, we take advantage of the value analysis's high performance with explicit values and only create boolean formulas when non-deterministic values occur. As a downside, symbolic execution, as it is based on value analysis, is far less potent in

	Bitvector theory	Integer theory
correct	833	815
false positives	52	93
false negatives	0	5
timeouts	1970	1944
errors	182	175
memory errors	1	1

Figure 6: Results of SymbolicExecutionCPA using bitvector theory with floats and integer theory on SV-COMP 2015 test sets.
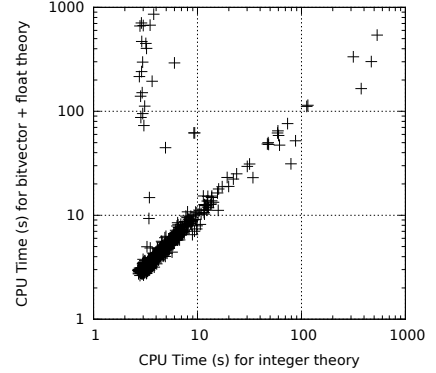
Figure 7: Runtime performance of symbolic execution using an integer theory in comparison to bitvector + float theory. Only benchmark runs that returned a result and for which the result of integer and bitvector was the same are displayed.

handling pointers than predicate analysis. But since this is only a very basic implementation, the potential of symbolic execution in the context of configurable software verification has still to be explored.

When choosing a SMT theory, the trade-off between performance and precision must be considered depending on the program's use of floats and bitwise operations. To decrease runtime without a loss of precision, the problem of path explosion should be tackled. A valid approach to this could be CEGAR, which is already implemented for the ValueAnalysisCPA. The introduction of compact symbolic execution as described in [8] could be another useful extension to allow more efficient handling of loops. By resolving the problem of path explosion, the number of iterations and as such also the number of SAT checks could be decreased, while keeping the high level of precision of symbolic execution. A preliminary implementation of CEGAR with symbolic execution already shows high potential. Refining this approach will be the main focus of future work.

5. REFERENCES

[1] BEYER, D. Second Competition on Software Verification (Summary of SV-COMP 2013). In *Proc. TACAS* (2013), vol. 7795 of *LNCS*, Springer, pp. 594–609.
[2] BEYER, D. Status Report on Software Verification (Competition Summary SV-COMP 2014). In *Proc. TACAS* (2014), vol. 8413 of *LNCS*, Springer,

pp. 373–388.

[3] Beyer, D. Software Verification and Verifiable Witnesses (Report on SV-COMP 2015). In *Proc. TACAS* (2015), LNCS, Springer. (To appear).

[4] Beyer, D. SV-COMP15: Benchmark Verification Tasks. `http://sv-comp.sosy-lab.org/2015/benchmarks.php`, 2015. Last checked: 03.31.2015.

[5] Beyer, D., Henzinger, T. a., and Théoduloz, G. Configurable Software Verification: Concretizing the Convergence of Model Checking and Program Analysis. In *Proc. CAV* (2007), vol. 4590 of *LNCS*, Springer, pp. 504–518.

[6] Beyer, D., and Keremoglu, M. E. CPAchecker: A tool for configurable software verification. In *Proc. CAV* (2011), vol. 6806 of *LNCS*, Springer, pp. 184–190.

[7] Beyer, D., and Löwe, S. Explicit-State Software Model Checking Based on CEGAR and Interpolation. In *Proc. FASE* (2013), vol. 7793 of *LNCS*, Springer, pp. 146–162.

[8] Slaby, J., Strejček, J., and Trtík, M. Compact Symbolic Execution. In *Proc. ATVA* (2013), LNCS, Springer, pp. 193–207.